

Linear time sorting, Heaps and HeapSort

1&1

1&1

Presentation by Antonio-Gabriel Sturzu

- Counting sort
- Radix sort
- Heaps
- Heap Sort
- Applications

- Assumes that each input of the array is in the interval $[0, k]$ and that $k = O(n)$
- Counts for every number how many elements are less than or equal to it
- Using this information it puts every number in its correct place
- It is a stable sort

Pseudocode:

COUNTING-SORT(A, B, k)

```
1  let  $C[0..k]$  be a new array
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $A.length$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  //  $C[i]$  now contains the number of elements equal to  $i$ .
7  for  $i = 1$  to  $k$ 
8       $C[i] = C[i] + C[i - 1]$ 
9  //  $C[i]$  now contains the number of elements less than or equal to  $i$ .
10 for  $j = A.length$  downto 1
11      $B[C[A[j]]] = A[j]$ 
12      $C[A[j]] = C[A[j]] - 1$ 
```

Counting sort

1&1

Example:

	1	2	3	4	5	6	7	8
A	2	5	3	0	2	3	0	3

	0	1	2	3	4	5
C	2	0	2	3	0	1

(a)

	0	1	2	3	4	5
C	2	2	4	7	7	8

(b)

	1	2	3	4	5	6	7	8
B							3	

	0	1	2	3	4	5
C	2	2	4	6	7	8

(c)

	1	2	3	4	5	6	7	8
B		0					3	

	0	1	2	3	4	5
C	1	2	4	6	7	8

(d)

	1	2	3	4	5	6	7	8
B		0				3	3	

	0	1	2	3	4	5
C	1	2	4	5	7	8

(e)

	1	2	3	4	5	6	7	8
B	0	0	2	2	3	3	3	5

(f)

- Used for sorting information keyed on multiple fields
- For example sorting some dates that are keyed on year, month and day
- But also it can be used to sort 64 bit integers in linear time
- Actually if we had infinite memory we could sort any vector of numbers in linear time no matter how big they are !

- Example for sorting 3 digit numbers:

329		720		720		329
457		355		329		355
657		436		436		436
839>>>>	457>>>>	839>>>>	457
436		657		355		657
720		329		457		720
355		839		657		839

- The pseudocode is extremely straightforward
- It is assumed that digit 1 is the least significant one and so one

RADIX-SORT(A, d)

1 **for** $i = 1$ **to** d

2 use a stable sort to sort array A on digit i

- If the chosen stable sort algorithm can sort in $\Theta(n+k)$ time where n is the number of elements and k is the maximum value for a digit then the total running time for radix sort will be $\Theta(d(n+k))$ where d is the maximum number of digits that a number in the input vector can have
- If the numbers in the vector don't have the same number of digits this is not a problem because we can pad them in order to achieve d digits

- We can prove that radix sort is correct by using induction on the current digit on which we must sort
- Induction: After k steps the numbers will be correctly sorted based on the first k digits
- Base case: After first step the numbers are correctly sorted after the first digit (the least significant digit)
- This is true because we are using a correct stable sort algorithm

- Now we assume $P(k)$ true and try to prove $P(k+1)$
- At the $k+1$ 'th step the sorting algorithm sorts the numbers on the $k+1$ 'th digit
- For different digits it is obvious that after $k+1$ steps the numbers are in correct order because the stable sorting algorithm will put the numbers with the lower value digits first
- It is not that obvious what happens if the digits are equal

- Now we use the fact that the sorting algorithm is stable and the induction step
- After k steps we know that the elements are correctly sorted
- So, for the elements with equal $k+1$ 'th digits if they preserve their relative order they will remain sorted
- Now, because the stable sort algorithm preserves the relative order for equal elements the proof is complete !

- On the general case:
 - If we have n b bit numbers and any positive integer $r \leq b$ then the running time of radix sort will be $\Theta((b/r)(n + 2^r))$ time if the stable sort it uses takes $\Theta(n + k)$ where the digits span in the $[0, k-1]$ interval
- The question that remains is when it is effective to use radix sort ?

- If $b=O(\log n)$ and we choose r approximately $\log n$ then the running time of radix sort will be $\Theta(n)$
- The problem is that radix sort uses much more memory than quicksort for example which sorts the vector in place if we don't take into account the stack usage
- Also the constant factor hidden by the Θ notation is bigger in practice than for quicksort
- Mainly it depends how well you implement the radix sort and if you want to optimize memory usage

- If you want to optimize for memory usage you might want to use an inplace sorting algorithm
- The choice also depends heavily on the type of input data that you have
- Problem:
 - Give an efficient algorithm for sorting a vector of 64 bit integers in linear time




- The binary heap represents an almost complete binary tree
- It comes in two flavours:
 - Max heap
 - Min heap
- In this presentation we will discuss only maxheaps
- But the same things presented for maxheaps also apply to minheaps as well

- A heap is usually stored as an array
- It uses the following three operations:
 - $\text{left}(i) = 2 \cdot i$
 - $\text{right}(i) = 2 \cdot i + 1$
 - $\text{parent}(i) = i / 2$
- A maxheap has the following property:
 - For every i $V[i] \leq V[\text{parent}(i)]$
- From this property we can conclude that the maximum element in a heap will be stored in the root

- Our first goal is how to build a heap out of an arbitrary array
- In order to do this we define a procedure that maintains the heap property called maxheapify
- The procedure receives as inputs the array and an index i
- It also assumes that the maxheap property is kept for the binary trees rooted at $\text{left}(i)$ and $\text{right}(i)$

- The problem might be that the tree rooted at node i does not obey the maxheap property
- This can happen if $A[i]$ is smaller than either of its children

- Example for maxheapify procedure:



- The pseudocode is straightforward:

```
MAX-HEAPIFY( $A, i$ )
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$ 
4       $\text{largest} = l$ 
5  else  $\text{largest} = i$ 
6  if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$ 
7       $\text{largest} = r$ 
8  if  $\text{largest} \neq i$ 
9      exchange  $A[i]$  with  $A[\text{largest}]$ 
10     MAX-HEAPIFY( $A, \text{largest}$ )
```





- The time complexity of maxheapify is $O(h)$ where h is the height of the tree rooted at i
- If the heap has n elements then h is $O(\log n)$
- The procedure is quite efficient except for the recursive call
- How could we eliminate the recursion ?

Heaps



Building the heap:

A

4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---



Building the heap:



- Pseudocode is quite simple:

BUILD-MAX-HEAP(*A*)

```
1  A.heap-size = A.length  
2  for i =  $\lfloor A.length/2 \rfloor$  downto 1  
3      MAX-HEAPIFY(A, i)
```

- Time complexity for building a maxheap
- A first obvious upper bound would be $O(n \log n)$ because we have n nodes, we call maxheapify on half of them and the maxheapify operation takes $O(\log n)$
- A tighter upper bound exists
- For a height h we have at most $\text{ceil}(n/2^{h+1})$ nodes

- The execution time will therefore be

$$\sum_{h=0}^{\lfloor \lg n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) = O \left(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h} \right)$$

- It turns out that this is actually $O(n)$!
- Now finally the heapsort algorithm
- Pseudocode:

HEAPSORT(A)

```
1  BUILD-MAX-HEAP( $A$ )
2  for  $i = A.length$  downto 2
3      exchange  $A[1]$  with  $A[i]$ 
4       $A.heap-size = A.heap-size - 1$ 
5      MAX-HEAPIFY( $A, 1$ )
```




Heapsort

1&1







Example:

A

1	2	3	4	7	8	9	10	14	16
---	---	---	---	---	---	---	----	----	----




Heapsort



Heapsort

1&1



- The time complexity is $O(n \log n)$
- The heapsort algorithm can be proved using an invariant
- Besides heapsort the main application of heaps is priority queues
- Also, heaps can be used to implement other tricky data structures where you need to retrieve the maximum and minimum in logarithmic time



Problems:

- How to efficiently merge k sorted arrays where the sum of the lengths of all the arrays is n ?
- How can a search engine keep in real time a list of the top k most frequent queries in that day ?
- Say we receive in real time a flow of n numbers that we don't know in advance. How can we efficiently maintain and retrieve the median values as we receive new numbers ?